

COMP0087 25/26

Lecture 4: Neural Language Models

23/01/2026

comp0087.github.io

Recap: N-gram Language Model

Goal: Estimate the probability of the next word given previous words.

$P(\text{UCL} \mid \text{I am studying AI at})$

Counting (“I am … at”, “I am … at UCL”)

Estimate probability by relative frequency

Key idea: Probability is estimated from observations in data

Recap: N-gram Language Model

Downsides:


- sparsity (lots of zeros)
- poor generalisation
- high cost (with the “5-grams at most”)

Solution: Neural Nets

N-gram modelling with neural nets

- Motivation: Address sparsity and poor generalisation in count-based n-gram models.
- Simple case: Feedforward Neural Network (FNN)
- Fixed-size context

N-gram modelling with neural nets



N-gram modelling with neural nets

I am studying NLP at University College London.

Self-supervised training data

< s > I am studying NLP at University College London < /s >

(< s > I am, studying)

(I am studying, NLP)

...

(at University College, London)

(University College London, < /s >)

How the input is constructed

(I am studying, NLP)

Look up in dictionary

(42, 15, 308) -> 127

Word	Index
I	42
am	15
NLP	127
...	...
studying	308

How the input is constructed

Look up embedding for each word

Each word index maps to a dense vector of dimension d
(e.g., $d = 128$)

Embedding Matrix E (size: $V \times d$, where V = vocabulary size)

Word "I" (index 42) $\rightarrow E[42] = [0.12, -0.45, 0.78, \dots]$

Word "am" (index 15) $\rightarrow E[15] = [-0.33, 0.21, 0.05, \dots]$

Word "studying" (index 308) $\rightarrow E[308] = [0.67, 0.14, -0.29, \dots]$

Embedding Matrix Initialisation

From pretrained word embedding (word2vec)

- Pros: Faster convergence, better generalisation.
- Cons: Domain mismatch.

Random initialised (uniform, gaussian)

How the input is constructed

Concatenate the 3 embeddings

Input vector $x = [E[42] ; E[15] ; E[308]]$

Final input dimension: $3 \times d$ (e.g., $3 \times 128 = 384$)

Hidden Layer Computation

Hidden layer transformation

$$\mathbf{h} = \text{activation}(\mathbf{W}_h \mathbf{x} + \mathbf{b}_h)$$

\mathbf{W}_h : hidden layer weight matrix (size: `hidden_size` × 3d)

\mathbf{b}_h : hidden layer bias vector (size: `hidden_size`)

Output Layer: Computing Word Probabilities

Project hidden state to vocabulary size

$$\mathbf{z} = \mathbf{W}_o \mathbf{h} + \mathbf{b}_o$$

\mathbf{W}_o : output weight matrix (size: $V \times \text{hidden_size}$)

\mathbf{b}_o : output bias vector (size: V)

\mathbf{z} : logits (raw scores for each word in vocabulary)

Output Layer: Computing Word Probabilities

Normalize logits to a valid distribution

$$\hat{\mathbf{y}} = \text{softmax}(\mathbf{z})$$

$$\hat{y}_i = P(w_i \mid \text{context}) = \frac{\exp(z_i)}{\sum_j^V \exp(z_j)}$$

Result: Probability distribution over entire vocabulary

Output Layer: Computing Word Probabilities

Word	Logit	Probability
NLP	2.8	0.35
hard	1.5	0.12
math	1.2	0.08
...		

Training: Loss computation

Ground truth (NLP): [0, 0, 0, ..., 1, 0, 0, 0]

Cross-Entropy Loss:

$$\text{Loss} = - \sum_i^V y_i \log(\hat{y}_i) = -\log(\hat{y}_{\text{target}})$$

Example: Loss = $-\log(0.35) = 1.05$

Training: gradient update

Trainable parameters $\theta = \{\mathbf{E}, \mathbf{W}_h, \mathbf{b}_h, \mathbf{W}_o, \mathbf{b}_o\}$

Update parameters: $\theta \leftarrow \theta - \alpha \frac{\partial L}{\partial \theta}$

α : learning rate (e.g., 0.001)

Summary

Step	Operation	Output
1	Word -> Index	42
2	Index -> Embedding	e_1, e_2, e_3
3	Concat	$x = [e_1; e_2; e_3]$
4	Hidden later	$h = \text{ReLU}(W_h x + b_h)$
5	Output layer	$z = W_o h + b_o$
6	Softmax	$\hat{y} = \text{softmax}(z)$

NN-LM Performance

Bonus

You get a learned word embedding for free.

Year: 2003

The quality of learned versus pretrained?

Downside of NN-LM

I put the beer in the fridge because it was warm.

Downside of NN-LM

- Cannot model long-range dependencies
- Same core limitation as n-gram models


Downside of NN-LM

Scaling issue: limitation of N


weight matrix grows rapidly

$$x \in \mathbb{R}^{nd_e} \quad W \in \mathbb{R}^{d_h \times nd_e}$$

Downside of NN-LM



$$x \in \mathbb{R}^{3d}$$




$$x \in \mathbb{R}^d$$

How can we dealing with variable length of input?


Idea

Use the same weight matrix to process information at different time steps?


A simple idea



Problem




Fix




Recurrent Neural Network

- Variable-length context: Can process sequences of any length
- Hidden state memory: Maintains information from all previous words
- Parameter sharing: Same weights applied at each time step

Recurrent Neural Network



Recurrent Neural Network



Recurrent Neural Network


- At each time step t , RNN takes:
- Current word embedding: \mathbf{x}_t
- Previous hidden state: \mathbf{h}_{t-1} (memory from past)
- And produces:
- New hidden state: \mathbf{h}_t (updated memory)

Recurrent Neural Network


$$h_t = \tanh(W_h h_{t-1} + W_x x_t + b)$$

- \mathbf{W}_h : hidden-to-hidden weight matrix (size: `hidden_size` × `hidden_size`)
- \mathbf{W}_x : input-to-hidden weight matrix (size: `hidden_size` × `embedding_dim`)
- \mathbf{b} : bias vector
- \tanh : activation function (keeps values between -1 and 1)

Recurrent Neural Network




Recurrent Neural Network



Recurrent Neural Network

$$h_t = \tanh(W_h h_{t-1} + W_x x_t + b)$$



<S>


|

am

UCL

Same weights
(W_h, W_x) for each
step.

RNN Language Model



RNN Language Model

After computing hidden state h_t , we can predict the next word

$$\mathbf{z} = \mathbf{W}_o \mathbf{h} + \mathbf{b}_o$$

\mathbf{W}_o : output weight matrix (size: $V \times \text{hidden_size}$)

\mathbf{b}_o : output bias vector (size: V)

\mathbf{z} : logits (raw scores for each word in vocabulary)

Output Layer: Computing Word Probabilities

Normalize logits to a valid distribution

$$\hat{\mathbf{y}} = \text{softmax}(\mathbf{z})$$

$$\hat{y}_i = P(w_i \mid \text{context}) = \frac{\exp(z_i)}{\sum_j^V \exp(z_j)}$$

Result: Probability distribution over entire vocabulary

Output Layer: Computing Word Probabilities

Time	Input	Hidden state	Prediction Target
1	I	h_1	am (0.42)
2	am	h_2	studying (0.35)
3	studying	h_3	NLP (0.28)
4	NLP	h_4	at (0.31)

$$h_3 \approx \alpha \text{Info("I")} + \beta \text{Info("am")} + \gamma \text{Info("studying")}$$

Training Objective

$$\text{Loss per step} = - \sum_i^V y_i \log(\hat{y}_i) = -\log(\hat{y}_{\text{target}})$$


$$\text{Loss} = -\frac{1}{T} \sum_t^T \log(\hat{y}_{t,\text{target}})$$

Training Objective

$$L_0 = -\log(\hat{y}["I"]) = -\log(0.42) = 0.87$$

$$L_1 = -\log(\hat{y}["am"]) = -\log(0.35) = 1.05$$

I (0.42)	am (0.35)	studying (0.5)	</s> (0.5)
-------------	--------------	-------------------	---------------



Training Objective


Backpropagation Through Time (BPTT):

- Gradients flow backward through time steps
- Update same weights (W_h , W_x , W_o) using accumulated gradients
- Challenge: Vanishing/exploding gradients over long sequences

RNN Language Model

- We isolate the shape of weight from the sequence length. Model weights (W_h , W_x) are independent of sequence length, making it possible to process arbitrary length of document.
- We can use the hidden state to store information that is lots of steps back, which makes long context modelling possible.

Decoder-only LM




RNN LM and GPT-2

- Autoregressive: Predict one token at a time, left-to-right
- Causal: Only use previous context, not future tokens
- Same objective: Minimize $-\log(\hat{y}_{\text{target}})$ at each position
- Only difference: recurrent versus self-attention

Limitation of RNN LM

I am a student
4 tokens


Je suis étudiant
3 tokens



Limitation of RNN LM

Modern solution:

< s > English: I am a student. French: Je suis étudiant < /s >



Translation Task

I am a student


Je suis étudiant

- Different languages, different input/output length
- Similar semantics for input and output

Sequence-to-sequence model

- Encoder RNN: Process input sequence, compress into context vector
- Decoder RNN: Generate output sequence from context vector

Sequence-to-sequence model



Sequence-to-sequence model

Encoder RNN (processes source sentence):

- At each time step t in the source:

$$h_t^{enc} = \tanh\left(W_h^{enc} h_{t-1}^{enc} + W_x^{enc} x_t + b^{enc}\right)$$

Context vector (final encoder state):

$$c = h_T^{enc}$$

T is the length of source sentence

Sequence-to-sequence model

Decoder RNN (generates target sentence):

Initialisation of decoder RNN state $s_0 = c$

At each time step t in the target:

$$s_t = \tanh(W_h^{dec} s_{t-1}^{dec} + W_x^{dec} y_{t-1} + b^{dec})$$


Output distribution:

$$P(y_t | y_1, \dots, y_{t-1}, \text{source}) = \text{softmax}(W_o \cdot s_t + b_o)$$

Sequence-to-sequence model

- Separate parameters: Encoder (W^{enc}) and Decoder (W^{dec}) have different weights
- Decoder input: Previous target word y_{t-1} (during training: ground truth; during inference: predicted word)
- Context c flows into decoder initialization or at each step

Training Seq2seq model



Training Seq2seq model

- Training (Teacher Forcing): At each decoder step, use the true previous word as input.

Step 1: Input <s> → $s_1 \rightarrow$ predict "Je" $\text{Loss}_1 = -\log P(\text{"Je"}|s_1)$

Step 2: Input "Je" → $s_2 \rightarrow$ predict "suis" $\text{Loss}_2 = -\log P(\text{"suis"}|s_2)$

...

Step 4: Input "étudiant" → $s_4 \rightarrow$ predict <EOS> $\text{Loss}_4 = -\log P(\text{<EOS>}| s_4)$


Training Seq2seq model

- Training (Teacher Forcing): At each decoder step, use the true previous word as input.

$$\text{Loss} = -\frac{1}{T} \sum_{t=1}^T \log P(y_t^* | y_1^*, \dots, y_{t-1}^*, \text{source})$$

where y_t^* is the true target word

Inference Seq2seq model



Model learns when to stop

- There is no predefined sequence length.
- Model generates a special <EOS> / stop token.
- Generation ends when stop token has highest probability.

Many NLP tasks are seq2seq

- Text summarisation

Original Text (truncated): lagos, nigeria (cnn) a day after winning nigeria's presidency, *muhammadu buhari* told cnn's christiane amanpour that **he plans to aggressively fight corruption that has long plagued nigeria** and go after the root of the nation's unrest. *buhari* said he'll "rapidly give attention" to curbing violence in the northeast part of nigeria, where the terrorist group boko haram operates. by cooperating with neighboring nations chad, cameroon and niger, **he said his administration is confident it will be able to thwart criminals** and others contributing to nigeria's instability. for the first time in nigeria's history, the opposition defeated the ruling party in democratic elections. *buhari* defeated incumbent goodluck jonathan by about 2 million votes, according to nigeria's independent national electoral commission. **the win comes after a long history of military rule, coups and botched attempts at democracy in africa's most populous nation.**

Baseline Seq2Seq + Attention: **UNK UNK** says his administration is confident it will be able to **destabilize nigeria's economy**. **UNK** says his administration is confident it will be able to thwart criminals and other **nigerians**. **he says the country has long nigeria and nigeria's economy.**

Seq2Seq Philosophy: Learning What's Important

Task: Text Summarization - Capturing Key Information

- Source document: Very long, rich information (hundreds of words)
- Summary: Short, only the most important points (tens of words)
- Question: How do we decide what's important?

Seq2Seq Philosophy: Learning What's Important

Classical method (Feature Engineering):

- TF-IDF: Extract sentences with high-scoring keywords
- Position-based: First/last sentences are important
- Keyword extraction: Look for specific indicator words

Learning from data

Encoder: Compress entire document into context vector c , context vector contains ALL information from source

Decoder: Learn to extract key information from c

- Trained on (document, summary) pairs
- Model learns what's important from the data
- No manual feature engineering needed!

Seq2seq to Seq-concat-seq

I am a student



Je suis étudiant

Modern solution:

< s > English: I am a student. French: Je suis étudiant < /s >

Learning from data

FineWeb Pretrain Data: Translation Pairs

Old English : bacan = "baker"

Proto-Germanic : bakana = "bake"

Old Norse : baka = "bake"

Middle Dutch : backen = "bake"

Old High German : bahhan = "bake"

German : backen = "bake"

PIE : bheg- "to warm, roast, bake"

Greek : phogein "to roast"

Root : bhe- "to warm"

From seq2seq to ChatGPT

- If we frame seq2seq as a language modeling task:
- Input + Output = one long sequence
- Train to predict next token (same as RNN-LM)
- No need for separate encoder and decoder
- This insight led to models like ChatGPT.

Thank you!