

COMP0087

Statistical Natural Language Processing

Lecture 11 – Reasoning

What is reasoning?

Logic: The process of deriving conclusions from premises according to valid rules of inference.

Cognitive Science: The cognitive process by which an agent draws inferences, forms judgments, or updates beliefs.

Symbolic AI: Algorithmic manipulation of symbolic representations to produce new representations consistent with a knowledge base.

Is Socrates mortal?

Socrates is a man.

Men are mortal.

Yes, Socrates is mortal.

$3x-1=2$. What is x ?

$(3x-1)+1=2+1$

$3x=3$

$3x/3=3/3$


$x=1$

A perspective: System 1 and System 2

System 1: thinking fast, typically associated with neural models

System 2: thinking slow, typically associated with symbolic reasoning

And for a long while the consensus was that we need neuro-symbolic solutions for difficult problems (i.e., problems that require System 2 type cognition)



A Proof-of-Concept – System 1 + System 2

Improving Coherence and Consistency in Neural Sequence Models with Dual-System, Neuro-Symbolic Reasoning

Maxwell Nye*
MIT


Michael Henry Tessler
MIT
DeepMind

Joshua B. Tenenbaum
MIT


Brenden M. Lake
NYU
Facebook AI Research

<https://arxiv.org/pdf/2107.02794.pdf>

A Proof-of-Concept – System 1 + System 2



A Proof-of-Concept – System 1 + System 2



Chain of Thought (CoT) in LLMs

Standard Prompting

Model Input

Q: Roger has 5 tennis balls. He buys 2 more cans of tennis balls. Each can has 3 tennis balls. How many tennis balls does he have now?

A: The answer is 11.

Q: The cafeteria had 23 apples. If they used 20 to make lunch and bought 6 more, how many apples do they have?

Model Output

A: The answer is 27. ❌

Chain-of-Thought Prompting


Model Input

Q: Roger has 5 tennis balls. He buys 2 more cans of tennis balls. Each can has 3 tennis balls. How many tennis balls does he have now?

A: Roger started with 5 balls. 2 cans of 3 tennis balls each is 6 tennis balls. $5 + 6 = 11$. The answer is 11.

Q: The cafeteria had 23 apples. If they used 20 to make lunch and bought 6 more, how many apples do they have?

A: The cafeteria had 23 apples originally. They used 20 to make lunch. So they had $23 - 20 = 3$. They bought 6 more apples, so they have $3 + 6 = 9$. The answer is 9. ✅



<https://arxiv.org/pdf/2201.11903>

Scratchpad

Show Your Work: Scratchpads for Intermediate Computation with Language Models

Maxwell Nye, Anders Johan Andreassen, Guy Gur-Ari, Henryk Michalewski, Jacob Austin, David Bieber, David Dohan, Aitor Lewkowycz, Maarten Bosma, David Luan, Charles Sutton, Augustus Odena

Input:

2 9 + 5 7

Target:

<scratch>

2 9 + 5 7 , C: 0

2 + 5 , 6 C: 1 # added 9 + 7 = 6 carry 1

, 8 6 C: 0 # added 2 + 5 + 1 = 8 carry 0

0 8 6

</scratch>

8 6

Scratchpad

Show Your Work: Scratchpads for Intermediate Computation with Language Models

Maxwell Nye, Anders Johan Andreassen, Guy Gur-Ari, Henryk Michalewski, Jacob Austin, David Bieber, David Dohan, Aitor Lewkowycz, Maarten Bosma, David Luan, Charles Sutton, Augustus Odena




Figure 3: Using a scratchpad significantly improves the performance of pre-trained Transformer-based models on addition, including their ability to generalize out of the training distribution to numbers with more digits. Models were trained on 1-8 digit addition. The baseline models were trained without intermediate scratchpad steps.


Data is still important – Minerva model

builds on general-purpose pretrained models by *further training them on a large, high-quality dataset of scientific and mathematical technical content.*

Mathematical Web Pages (~17.5 B tokens)


arXiv Papers (~21 B tokens)

Accuracy on MATH Subtopics (Base PaLM vs Minerva)



July 25, 2024 Science

AI achieves silver-medal standard solving International Mathematical Olympiad problems



Process infographic of AlphaProof's reinforcement learning training loop: Around one million informal math problems are translated into a formal math language by a formalizer network. Then a solver network searches for proofs or disproofs of the problems, progressively training itself via the AlphaZero algorithm to solve more challenging problems.

AlphaProof: a formal approach to reasoning

AlphaProof is a system that trains itself to prove mathematical statements in the formal language Lean. It couples a pre-trained language model with the AlphaZero reinforcement learning algorithm, which previously taught itself how to master the games of chess, shogi and Go.

Formal languages offer the critical advantage that proofs involving mathematical reasoning can be formally verified for correctness. Their use in machine learning has, however, previously been constrained by the very limited amount of human-written data available.

In contrast, natural language based approaches can hallucinate plausible but incorrect intermediate reasoning steps and solutions, despite having access to orders of magnitudes more data. We established a bridge between these two complementary spheres by fine-tuning a Gemini model to automatically translate natural language problem statements into formal statements, creating a large library of formal problems of varying difficulty.

July 21, 2025 Research

Advanced version of Gemini with Deep Think officially achieves gold-medal standard at the International Mathematical Olympiad

While our approach this year was based purely on natural language with Gemini, we also continue making progress on our formal systems, AlphaGeometry and AlphaProof. We believe agents that combine natural language fluency with rigorous reasoning - including verified reasoning in formal languages - will become invaluable tools for mathematicians, scientists, engineers, and researchers, helping us advance human knowledge on the path to AGI.



Andrew Lampinen
@AndrewLampinen



Our LM (and others) achieved gold-medal performance in the international Math Olympiad ([x.com/demishassabis/...](https://x.com/demishassabis/)). Importantly, unlike last year's silver, this system used natural language end-to-end, rather than relying on a verifiable formal language.



Demis Hassabis ✓ @demishassabis · Jul 21, 2025

Official results are in - Gemini achieved gold-medal level in the International Mathematical Olympiad! 🏆 An advanced version was able to solve 5 out of 6 problems. Incredible progress - huge congrats to @lmthang and the team! [deepmind.google/discover/blog/...](https://deepmind.google/discover/blog/)

10:46 PM · Jul 21, 2025 · 11.9K Views


<https://deepmind.google/blog/advanced-version-of-gemini-with-deep-think-officially-achieves-gold-medal-standard-at-the-international-mathematical-olympiad/>

System 2




A talk by Yoshua Bengio on System 1 and 2: <https://www.youtube.com/watch?v=FtUbMG3rlFs>
And a “debate” between him and Gary Marcus: <https://www.youtube.com/watch?v=EeqwFigFvJA>


Self-Taught Reasoner (STaR) - Basic version




CommonsenseQA with STaR (Basic)




What to do with the wrong answers?



CommonsenseQA with STaR




Imitation Learning works too





Prompt LLMs like GPT-4 to generate reasoning trajectories and use this collected data to fine-tune another, smaller, or weaker model

Collecting diverse solutions for the same problems



Fixing failed trajectories using teacher model





Using RL to teach model to reason:

The Deepseek R1

Deepseek R1 (Reasoning LLMs)


DeepSeek-R1: Incentivizing Reasoning Capability in LLMs via Reinforcement Learning

DeepSeek-AI

research@deepseek.com

Abstract

We introduce our first-generation reasoning models, DeepSeek-R1-Zero and DeepSeek-R1. DeepSeek-R1-Zero, a model trained via large-scale reinforcement learning (RL) without supervised fine-tuning (SFT) as a preliminary step, demonstrates remarkable reasoning capabilities. Through RL, DeepSeek-R1-Zero naturally emerges with numerous powerful and intriguing reasoning behaviors. However, it encounters challenges such as poor readability, and language mixing. To address these issues and further enhance reasoning performance, we introduce DeepSeek-R1, which incorporates multi-stage training and cold-start data before RL. DeepSeek-R1 achieves performance comparable to OpenAI-o1-1217 on reasoning tasks. To support the research community, we open-source DeepSeek-R1-Zero, DeepSeek-R1, and six dense models (1.5B, 7B, 8B, 14B, 32B, 70B) distilled from DeepSeek-R1 based on Qwen and Llama.



What is the reward in Deepseek-R1-Zero? Bunch of rules

The reward is the source of the training signal, which decides the optimization direction of RL. To train DeepSeek-R1-Zero, we adopt a rule-based reward system that mainly consists of two types of rewards:

- **Accuracy rewards:** The accuracy reward model evaluates whether the response is correct. For example, in the case of math problems with deterministic results, the model is required to provide the final answer in a specified format (e.g., within a box), enabling reliable rule-based verification of correctness. Similarly, for LeetCode problems, a compiler can be used to generate feedback based on predefined test cases.
- **Format rewards:** In addition to the accuracy reward model, we employ a format reward model that enforces the model to put its thinking process between '<think>' and '</think>' tags.

A conversation between User and Assistant. The user asks a question, and the Assistant solves it. The assistant first thinks about the reasoning process in the mind and then provides the user with the answer. The reasoning process and answer are enclosed within <think> </think> and <answer> </answer> tags, respectively, i.e., <think> reasoning process here </think> <answer> answer here </answer>. User: **prompt**. Assistant:

Table 1 | Template for DeepSeek-R1-Zero. **prompt** will be replaced with the specific reasoning question during training.

The aha moment!

Question: If $a > 1$, then the sum of the real solutions of $\sqrt{a - \sqrt{a + x}} = x$ is equal to

Response: <think>

To solve the equation $\sqrt{a - \sqrt{a + x}} = x$, let's start by squaring both $\dots$

$$\left(\sqrt{a - \sqrt{a + x}}\right)^2 = x^2 \implies a - \sqrt{a + x} = x^2.$$

Rearrange to isolate the inner square root term:

$$(a - x^2)^2 = a + x \implies a^2 - 2ax^2 + (x^2)^2 = a + x \implies x^4 - 2ax^2 - x + (a^2 - a) = 0$$

$\dots$

Wait, wait. Wait. That's an aha moment I can flag here.

Let's reevaluate this step-by-step to identify if the correct sum can be $\dots$

We started with the equation:

$$\sqrt{a - \sqrt{a + x}} = x$$


First, let's square both sides:

$$a - \sqrt{a + x} = x^2 \implies \sqrt{a + x} = a - x^2$$

Next, I could square both sides again, treating the equation: $\dots$

$\dots$

The multi-stage pipeline of DeepSeek-R1




R1-Zero vs O1

Model	AIME 2024		MATH-500	GPQA Diamond	LiveCode Bench	CodeForces
	pass@1	cons@64	pass@1	pass@1	pass@1	rating
OpenAI-o1-mini	63.6	80.0	90.0	60.0	53.8	1820
OpenAI-o1-0912	74.4	83.3	94.8	77.3	63.4	1843
DeepSeek-R1-Zero	71.0	86.7	95.9	73.3	50.0	1444

The basic Zero-R1 is competitive

“The self-evolution of DeepSeek-R1-Zero exemplifies how RL can autonomously enhance a model’s reasoning capabilities.”


28



Smaller LMs: RL from scratch vs. Distillation

Model	AIME 2024		MATH-500
	pass@1	cons@64	pass@1
QwQ-32B-Preview	50.0	60.0	90.6
DeepSeek-R1-Zero-Qwen-32B	47.0	60.0	91.6
DeepSeek-R1-Distill-Qwen-32B	72.6	83.3	94.3
DeepSeek-R1-Distill-Qwen-1.5B	28.9	52.7	83.9
DeepSeek-R1-Distill-Qwen-7B	55.5	83.3	92.8
DeepSeek-R1-Distill-Qwen-14B	69.7	80.0	93.9
DeepSeek-R1-Distill-Qwen-32B	72.6	83.3	94.3
DeepSeek-R1-Distill-Llama-8B	50.4	80.0	89.1
DeepSeek-R1-Distill-Llama-70B	70.0	86.7	94.5


Distillation from Stronger Reasoners





<https://github.com/open-thoughts/open-thoughts>

Thinking mode ...

Think Before Talking



Think Before Talking



Inference Time Scaling

Generate many plausible responses by default

Verification (will cover these in the next lecture)



Verify those responses and return the best one

Scaling Test-Time Compute

Scaling LLM Test-Time Compute Optimally can be More Effective than Scaling Model Parameters

Charlie Snell^{*,1}, Jaehoon Lee², Kelvin Xu^{*,2} and Aviral Kumar^{*,2}

^{*}Equal advising, ¹UC Berkeley, ²Google DeepMind, ^{*}Work done during an internship at Google DeepMind





Inference Time Scaling (emergence of OpenAI O1)

September 12, 2024 Release

Learning to reason with LLMs


We are introducing OpenAI o1, a new large language model trained with reinforcement learning to perform complex reasoning. o1 thinks before it answers—it can produce a long internal chain of thought before responding to the user.



“Thinking” mode

Instead of producing a short reasoning chain, they:

- Sample multiple reasoning trajectories
- Score or compare them
- Choose the most consistent
- Allocate more tokens to “thinking”
- Delay final answer emission




Language models show human-like content effects on reasoning

Ishita Dasgupta^{*,1}, Andrew K. Lampinen^{*,1}, Stephanie C. Y. Chan¹,
James L. McClelland^{1,2} and Felix Hill¹

^{*}Equal contributions, listed alphabetically, ¹DeepMind, ²Stanford University

<https://arxiv.org/pdf/2207.07051>

“Humans reason more reliably when the semantic content of a problem supports the correct logical Inferences”



Consistent

All **guns** are **weapons**.
All **weapons** are **dangerous things**.
All **guns** are **dangerous things**.

Violate

All **dangerous things** are **weapons**.
All **weapons** are **guns**.
All **dangerous things** are **guns**.

Nonsense

All **zoct** are **spuff**.
All **spuff** are **thrund**.
All **zoct** are **thrund**.

LLMs conflate their internal knowledge with given task knowledge Sometimes you want the model to only reason with given facts ... Sometimes you need guarantees that this actually happens.

Symbolic/Formal Reasoning is a way to achieve this.

Symbolic Representation

- Representing the meaning of natural language is ultimately a difficult philosophical question
- Traditional approach is to map ambiguous NL to unambiguous logic in **first-order logic (FOL)**, also known as predicate logic.
- Standard inference (theorem proving) methods exist for FOL that can determine when one statement entails (implies) another. Questions can be answered by determining what potential responses are entailed by given NL statements and background knowledge all encoded in FOL.

Autoformalization is the process of automatically translating from natural language to formal specifications and proofs.

Recognising Textual Entailment

Example: 78 (FALSE)

T: Clinton's new book is not big seller here.

H: Clinton's book is a big seller.

<https://aclanthology.org/H05-1079.pdf>

1

Tokenisation & Lemmatisation · CCG Parser

- Both T and H tokenised; lemmas extracted; wide-coverage CCG parser applied
- Handles complex syntax (apposition, coordination, negation) robustly

2

Semantic Parsing → DRS Construction · Discourse Representation Theory

- CCG parses converted to Discourse Representation Structures (DRS)
- Neo-Davidsonian event semantics; proper names resolved anaphorically
- Scope fully specified; negation, disjunction handled recursively

3

DRS → First-Order Logic Translation · Kamp & Reyle (1993)

- Standard FOL translation of DRS(T) and DRS(H) ready for automated reasoners
- e.g. $\text{soar}(e) \wedge \text{agent}(e, \text{prices}) \rightarrow \text{FOL}(\text{DRS}(T))$

4

Background Knowledge Injection · WordNet · CIA Factbook · Manual Axioms

- Lexical: WordNet hyponymy $\rightarrow \forall x(\text{soar}(x) \rightarrow \text{rise}(x))$; synset-sisters $\rightarrow$ negation axioms
- Generic: ~20 hand-crafted axioms (active-passive, spatial containment, possessives)
- Geographic: auto-extracted $\forall x \forall y (\text{paris}(x) \wedge \text{france}(y) \rightarrow \text{in}(x, y))$
- WSD via Lesk algorithm maps predicates to WordNet synsets

SHALLOW

Word Overlap (wnoverlap)

Per H lemma: check equality, WordNet synonymy or derivation vs T; weight by IDF (Google API). Score normalised $\in [0, 1]$.

Length Features

$\text{len}(T)$, $\text{len}(H)$, $\text{len}(H)/\text{len}(T)$ — true entailments tend to have shorter H.

DEEP SEMANTIC

Theorem Prover (Vampire)

Proves $\text{FOL}(T) \rightarrow \text{FOL}(H)$ or inconsistency. High precision, very low recall (~30/800 proofs).

Model Builder (Paradox)

Seeks counterexample for $\neg A$. Domain/model size diff approximates entailment when proof unavailable.

Same but now with LLMs

LLM

No giant language model could have bad performance.
If a language model has good performance, it is used by some researchers.
A work used by some researchers should be popular.
If BERT is a giant language model, then the same for GPT3.
BERT is a giant language model.


Is the following statement true, false, or unknown? GPT3 is popular.



Facts:

- $\neg(\exists x(\text{LanguageModel}(x) \wedge \text{Giant}(x) \wedge \neg\text{GoodPerformance}(x)))$
- $\forall x(\text{LanguageModel}(x) \wedge \text{GoodPerformance}(x) \rightarrow \text{UsedbySomeReseachers}(x))$
- $\forall x(\text{UsedbySomeResearchers}(x) \rightarrow \text{Popular}(x))$
- $\text{LanguageModel}(\text{bert}) \wedge \text{Giant}(\text{bert}) \rightarrow \text{LanguageModel}(\text{gpt3}) \wedge \text{Giant}(\text{gpt3})$
- $\text{Language}(\text{bert})$
- $\text{Giant}(\text{bert})$

Query: Popular(gpt3)



Entailment

<https://arxiv.org/abs/2305.12295>

Dataset	ChatGPT (gpt-3.5-turbo)			GPT-3.5 (text-davinci-003)			GPT-4 (gpt-4)		
	Standard	CoT	Logic-LM	Standard	CoT	Logic-LM	Standard	CoT	Logic-LM
PrOntoQA	47.40	<u>67.80</u>	61.00	51.80	83.00	<u>85.00</u>	77.40	<u>98.79</u>	83.20
ProofWriter	35.50	49.17	<u>58.33</u>	36.16	48.33	<u>71.45</u>	52.67	68.11	<u>79.66</u>
FOLIO	45.09	57.35	<u>62.74</u>	54.60	57.84	<u>61.27</u>	69.11	70.58	<u>78.92</u>

Can agents reason in the wild?

- LLMs are good at solving well-defined reasoning problems, **but real-world scenarios are complex and challenging**
- We introduce the task of **reasoning in the wild**: LLMs to solve problems of hybrid types by decomposing them into subproblems and applying corresponding formalisms/tools to solve them step by step.
- Created problems that are (1) ambiguous; and (2) requires multiple tactics to solve
 - Take existing problems (e.g., GSM8K, FOLIO, ReClor) and “blend” them into a hybrid problem

[Can LLMs Reason in the Wild with Programs? Yang, etl. 2024](#)

Ambiguous, Mixed-in-scope Problem

Context:

Jerry works at Acme Corporation, whose strategic decisions are set to significantly affect its employees' daily routines. Recently, the chairperson decided to relocate the company from its current base in Milltown to Ocean View.

...

To celebrate the relocation, Jerry helped preparing candy gift bags. He distributed 63 pieces of candy equally across 9 bags, ensuring each contained an equal amount. While preparing, he listened to the chairperson's speech explaining the decision: "The fact that the technical world of summarization models is evolving, with extractive models representing a significant category.

...

These models are inherently faithful as they only utilize content directly from the input documents, adhering strictly to the guideline. That's why we need to move", said the chairperson.

...



Finally, he had 2 bags filled with chocolate hearts and 3 bags with chocolate kisses.

Which one of the following is correct:

1. The number of candies Jerry put that were not chocolate is 26.
2. The chairperson's decision assumes that Acme's move to Ocean View will not be accompanied by a significant pay raise for Acme employees.
3. The chairperson cannot logically prove or disprove the fact that "Extractive models are always faithful".

Reasoning in the wild – Inference Flow


- **Tactic:** a specific formalism defined for a problem type
- **Routing:** given a problem of unknown type, the agent chooses the right tactic to approach it
- **Execution:** the agent performs T thought, A action, and receives O observation at every round
- **Iteration:** the agent iteratively identifies the subproblems, solve them, and produce the final answer



example of FOL tactic

Experiments – Standalone vs. Hybrid

Performance
on Standalone
Problems



Performance
on Hybrid
Problems

