

COMP0087 25/26

Lecture 1: Introduction

13/01/2026

comp0087.cs.ucl.ac.uk

Who are we?

Instructors

Yao Lu

Lecturer in Natural Language Processing

Contact: yao.lu@cs.ucl.ac.uk

Research: language model pretraining, prompt optimisation

Instructors

Ehsan Shareghi

Associate Professor in Natural Language Processing

Contact: ehsan.shareghi@gmail.com

Research: LLM safety, speech processing

Instructors

Pontus Stenetorp

Professor in Natural Language Processing

Contact: p.stenetorp@cs.ucl.ac.uk

Research: Anything that fits into NLP (Multilinguality, data, knowledge, generalisation, etc)

Instructors

Vasileios (Bill) Lampos

Associate Professor

Contact: v.lampos@ucl.ac.uk

Research: Sequence modelling (time series forecasting), LLMs training dynamics (e.g., mode collapse)

Assumptions about you

You already know the basics of applied ML:

- What the training/dev/test cycle is
- What a loss function is
- What an optimization technique does (at a high level)
- How to evaluate a model (given a standard benchmark)

Assumptions about you

You can already do these things:

- Perform basic data manipulation
- Build basic knowledge in Python

Who are you?

Level of experience in NLP

1. Beginner (no prior NLP experience)
2. Intermediate (coursework or projects)
3. Advanced (research or professional experience)

Who are you?

I'm taking this course to

1. Prepare for job searching
2. Learn how to conduct NLP research
3. Apply NLP to ongoing projects and research
4. Other

Module communication

General, non-personal, queries

Largest postgraduate module

General, non-personal, queries:

Please, pretty please use the Moodle forum

Personal queries

Personal queries:

comp0087@cs.ucl.ac.uk

General project-related queries

General, project-related queries:

Moodle forum

Your project-related queries

Specific, project-related queries:

Assigned teaching assistant

Administrative-related queries

Administrative logistics (registration, etc.):

cs.pgt-students@ucl.ac.uk

Module content

Learning goals

1. Foundational “building blocks” for LLMs
2. Advanced LLM development/applications
3. Practical, empirical understanding of building NLP systems
 - Analysis, empirical experiments, etc.
 - Group project

Teaching plan

Basic NLPs 3-4 weeks

Advanced NLPs (LLMs and applications) 3-4 weeks

Topic Lectures from Industry speakers (e.g., Meta AI, Microsoft etc) 2-3 weeks

NLP Course Projects (supervised by our awesome TAs and faculty members) 1-10 weeks

Teaching plan: Foundations

Tokenisation

What is your input?

Token
representation

In what form?

Teaching plan: Foundations

Traditional LM

Neural LM

How to model the data?

Teaching plan: Foundations

Attention

Transformer

Which architecture and why?

Teaching plan: Foundations

Prompting

How to "talk" to the model?

Teaching plan: Advanced Topics

Advanced NLPs

- Tuning Large Language Models
- Alignment with Human Feedback
- Reasoning LLMs
- Retrieval Augmented Generation (RAG)
- LLMs Pretraining

...

Teaching plan: Industry Lectures

Please let us know which topics you're interested in.

Group project

Assessment

Coursework (Group): 100% of module mark

- Mandatory bi-weekly progress meetings
- Single-page progress reports (due date TBD)
- Eight-page final report
- Not including references

Group formation

- Groups of *five* or *six*
- * Exceptions will be rare
- Group formation due 18 January by midnight
- * Start looking *now*
- * Decide on group name
- * Post UCL e-mails in dedicated Moodle forum

Project requirements

- Must involve language
- Empirical investigation
- Designed by students, with the *help* of us
- * Ultimate responsibility and control is with *you*
- Outcomes *may* be publishable
- * But *no* guarantees

Project timeline sketch

- Week 1 to 2: Project design and literature review
- Week 3 to 4: Data collection, experimental design, initial implementation
- Week 5 to 6: Experiments not working out: Revisions, revisions, revisions,
- Week 7 to 8: Additional experiments and analysis
- Week 9 to 10: Report writing

Project design

- Minimum viable project
- Align with group interests and teaching assistant expertise
- Design to be "split"

Computational resources

- *Limited*: Design your projects with this in mind!
- - UCL Computer Science:
- * <https://tsg.cs.ucl.ac.uk/gpus>
- * ~30 x NVIDIA RTX 4060 cards (16GB)
- * ~25 x NVIDIA RTX 3090 cards (24GB)

Computational resources

- Google Colab:
- * <https://colab.research.google.com>
- * Web "notebooks" with free GPUs and TPUs (LLM access?)
- * Can pay for better service (GPU server rental)
- co:here credits
- * Application process announced shortly

Project design example

I want to train a language model.

I want to train a language model that better
understands twitter posts.

Project design example

- What's broken? (Problem)
- How will you know it's fixed? (Metrics)
- What will you try? (Approach)
- Can you actually do it? (Feasibility)
- What if it fails? (Contingency)

Project design example

What's Broken?

- Why do current LLMs struggle with tweets?
- What evidence supports this? (Literature review)
- Show examples of failures

Project design example

How Will You Know It's Fixed?

- Which task? (sentiment, NER, sarcasm detection, etc.)
- Which dataset? (TweetEval? SemEval? Custom?)
- What's the baseline? (GPT-4? RoBERTa-Twitter? Prompting vs. fine-tuning?)

Project design example

What Will You Try?

- Which base model?
- Prompting, fine-tuning, or pretraining? (Keep budget in mind)
- What training data? (Source, size, licensing)
- Why should this approach work? (Hypothesis)

Project design example

Can You Actually Do It?

- GPU hours (check maximum available resources)
- Training data (size, cost)
- Evaluation data (quality, cost)
- Timeline (10 weeks in total)
- Split tasks

Project design example

What If It Fails?

- No improvement: Why? Data quality? Wrong task?
Write negative result analysis
- Reframe as benchmark or dataset contribution
- Discuss with your TA in advance

Teaching assistants

Adam Oomerjee-Vawda (adam.vawda@gmail.com)

- PhD student at UCL AI centre
- Interests: Efficient models, reasoning
- Favourite paper:
- Bottlenecked Transformers: Periodic KV Cache Abstraction for Generalised Reasoning. <https://arxiv.org/abs/2505.16950>

Edan Toledo (edan.toledo.24@ucl.ac.uk)

- PhD student at UCL and Facebook
- Interests:
 - * General search methods * Agentic systems * Meta-learning
 - * Reinforcement learning
- Favourite paper:
 - * "A general reinforcement learning algorithm that masters chess, shogi, and Go through self-play" by Silver et al. (2018)
 - * <https://www.science.org/doi/10.1126/science.aarr6404>

Eduardo Sánchez (eduardo.sanchez.22@ucl.ac.uk)

- final-year PhD student at UCL and Facebook
- Interests:
- * Low-resource languages * Multilinguality
- Favourite paper:
- * "Linguini: A benchmark for language-agnostic linguistic reasoning" by Sánchez et al. (2024) <https://arxiv.org/abs/2409.12126>

Jiayi Wang (jiayi.lin.wang@ucl.ac.uk)

- final-year PhD student
- Interests:
 - * Inclusive and efficient multilingual large language models
- Favourite paper:
 - * "Multilingual Pretraining Using a Large Corpus Machine-Translated from a Single Source Language" by Wang et al. (2023)
 - * <https://arxiv.org/abs/2410.23956>

Karen Hambardzumyan (karen.hambardzumyan.23@ucl.ac.uk)

- PhD student at UCL and Facebook
- Interests:
 - * Mechanistic interpretability * Multi-agent systems
- Favourite paper:
 - * "WARP: Word-level Adversarial ReProgramming" by Hambardzumyan et al. (2021) * <https://aclanthology.org/2021.acl-long.381>

Lovish Madaan (lovish.madaan.23@ucl.ac.uk)

- PhD student at UCL and Facebook (GenAI)
- Interests:
 - * Generalisation behaviours in reinforcement learning *
 - Better/harder evaluation
- Favourite paper:
 - * "Amortizing intractable inference in large language models" by Hu et al. (2024) * <https://arxiv.org/abs/2310.04363>

Ralph Tang (r-tang.25@ucl.ac.uk)

- PhD student at NLP group
- Interests: * Multimodality * LLM/VLM interpretability
- Favourite paper:
- What the DAAM: Interpreting Stable Diffusion Using Cross Attention. ACL 2023

Hossein A (Saeed) Rahmani
(hossein.rahamani.22@ucl.ac.uk)

- PhD student at NLP group
- Interests: * Evaluations and benchmarks * Self-Improvement
- Favourite paper:
- "Large Language Models Cannot Self-Correct Reasoning Yet" by Huang et al. (2024) <https://arxiv.org/abs/2310.01798>

Tokenisation

Let's study NLP together.



Human (Word-level)

Let's

study

NLP

together

.

Total: 5 tokens (words and punctuation)



GPT3.5 (Subword-level Tokenization)

Let

' s

study

N

LP

together

.

Total: 7 tokens (subword units for better vocabulary coverage)

From Text to Tokens

Tokens	Characters
7	25

Let's study NLP together.

[10267, 596, 4007, 452, 12852, 3871, 13]

Text

Token IDs

Text

Token IDs

The Reverse Spelling Problem

strawberry -> y-r-r-e-b-w-a-r-t-s

Why is this difficult?

The Reverse Spelling Problem

strawberry -> [“str”, “aw”, “berry”]

-> y-r-r-e-b-w-a-r-t-s?

Why is this difficult?

The Reverse Spelling Problem

strawberry -> [“str”, “aw”, “berry”]
-> [“yr”, “reb”, “warts”]

Why is this difficult?

The Reverse Spelling Problem

strawberry -> [496, 675, 15717]
-> [11160, 32575, 64156]

Why is this difficult?

The Number Comparison Problem

"Is 9.8 greater than 9.11?"

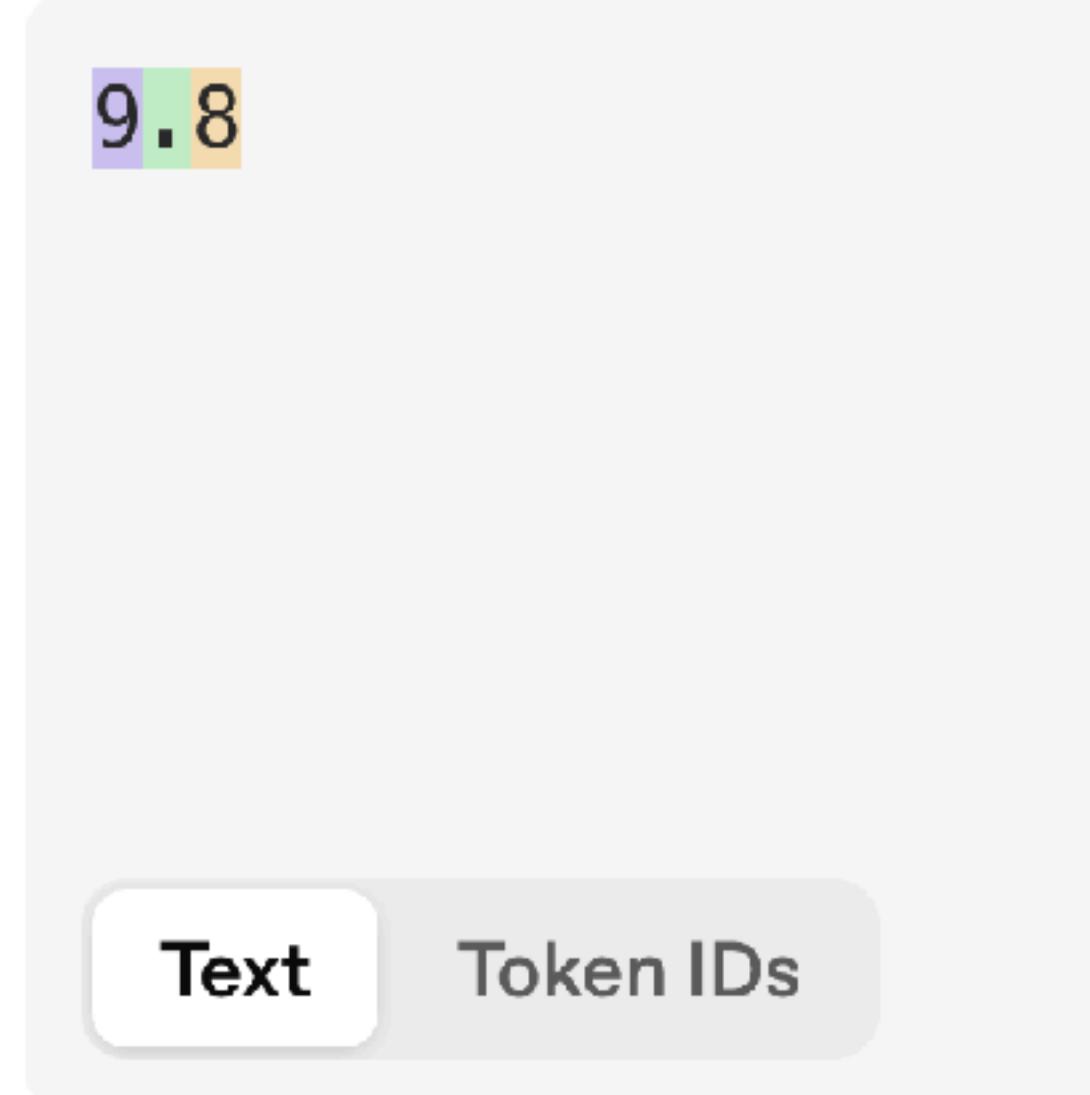
The Number Comparison Problem

Human: 9.**8** > 9.**1**1

The Number Comparison Problem

GPT4: 9.8 < 9.11

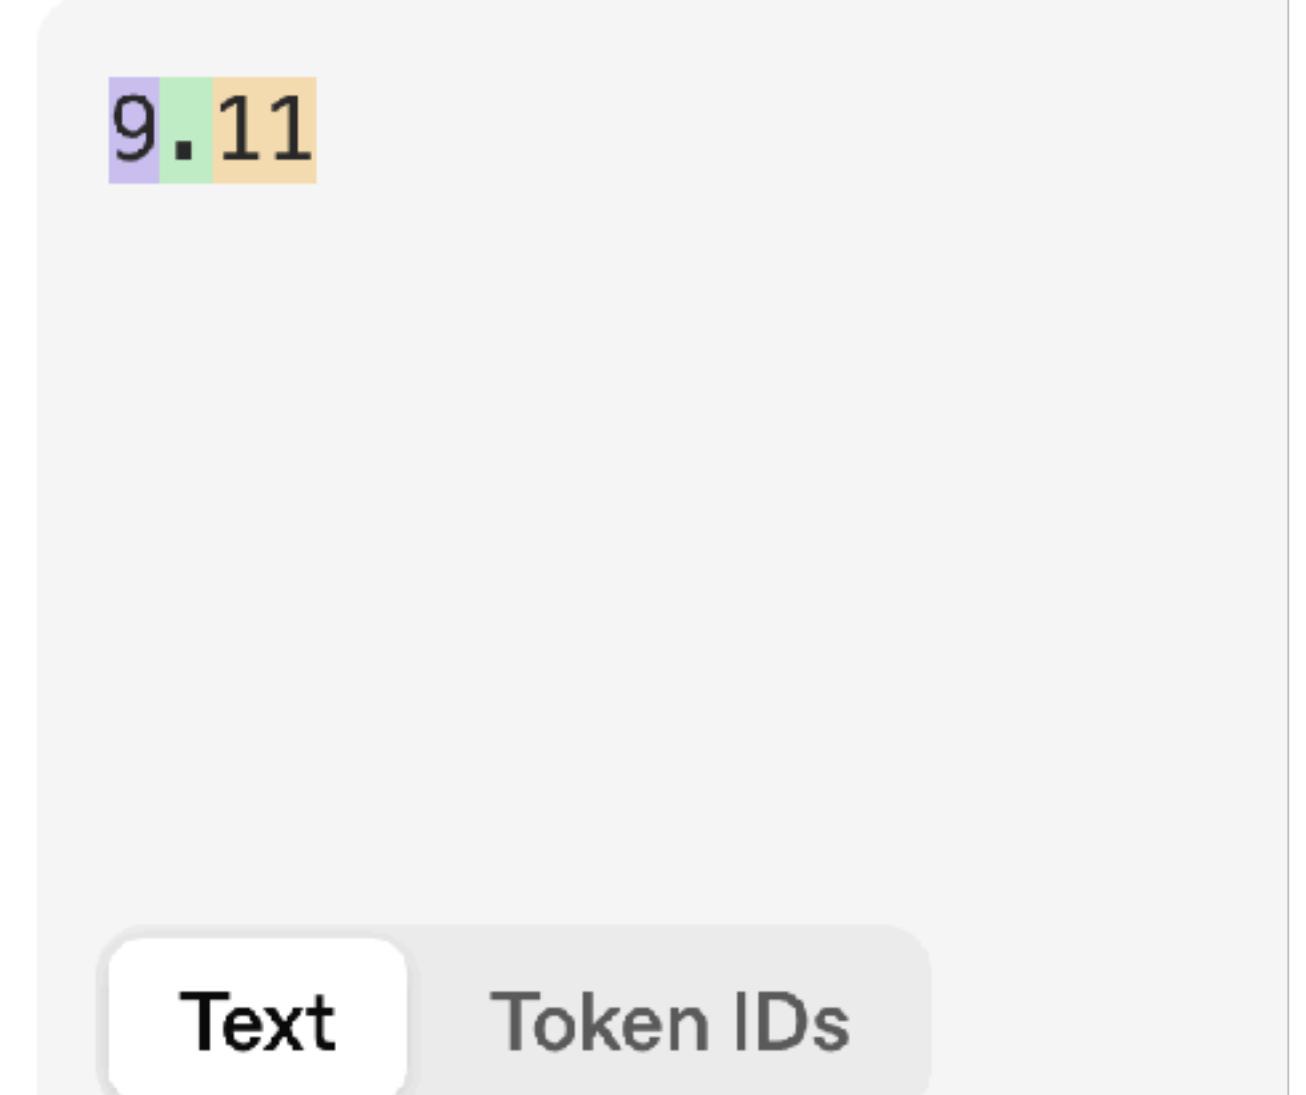
Tokens	Characters
3	3



9.8

Text Token IDs

Tokens	Characters
3	4



9.11

Text Token IDs

Definition of tokenisation

- The process of splitting **text** into **meaningful tokens**, which are then mapped to **numerical IDs**.

Classical tokenisation methods

Character-level tokenisation: strawberry -> [s, t, r, a, w, b, e, r, r, y]

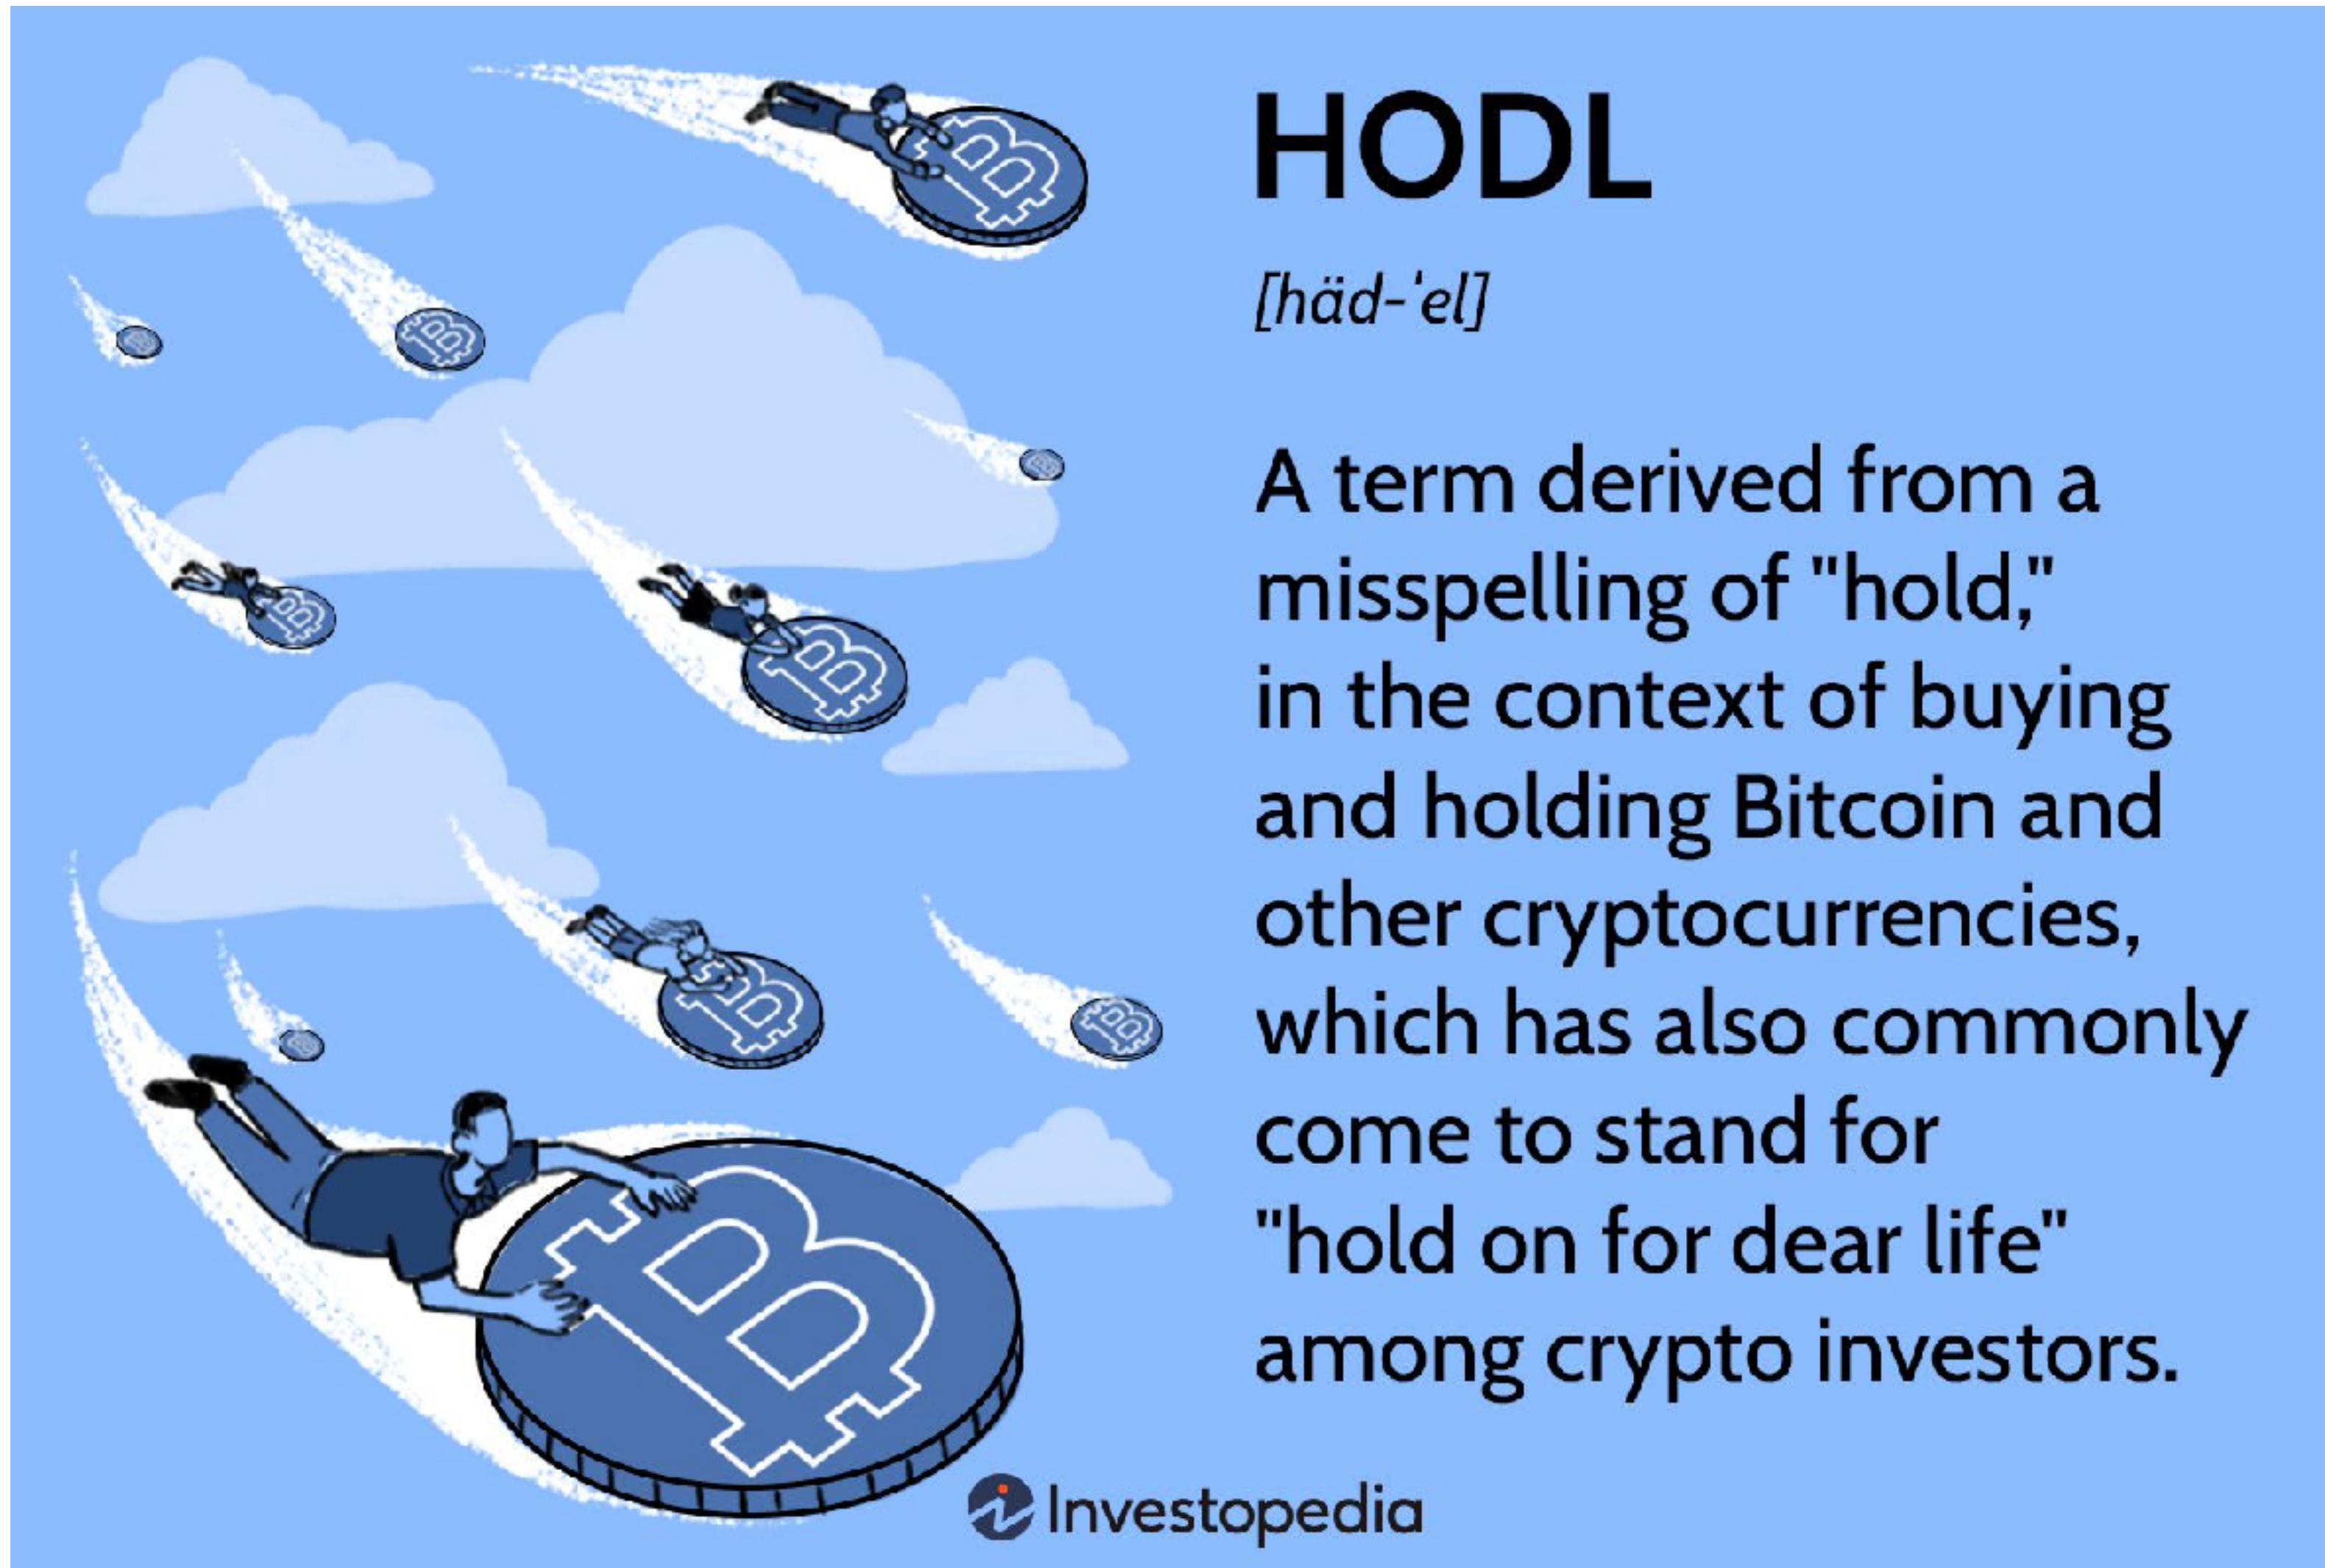
- Pros: a very small set of symbols (100?) can represent almost everything
- Cons: less meaningful segmentation

Classical tokenisation methods

Word-level tokenisation: “I am happy” -> [Index(I), Index(am), Index(happy)]

- Pros: meaningful and natural segmentation (by whitespace)
- Cons: huge vocabulary space

Out of Vocabulary Tokens



Out of Vocabulary Tokens

Solution: replace out of vocabulary tokens with <UNK>

“My favourite food is <UNK>”

“My favourite <UNK> is <UNK>”

“<UNK> favourite <UNK> is <UNK>”

Tokenisation design

	Character-based method	Word-level method	?
Vocabulary size	~100	100,000	1k-100k
Corpus coverage	High (any text)	Medium (OOV issues)	High
Meaningfulness	Low	High	High
Sequence length	Very long	Short	Medium

Subword Tokenisation

- Frequent words stay as single tokens ("cat").
- Rare/complex words split into meaningful sub-units ("tokenisation" -> "token" + "isation").
- Handles OOV: Can decompose new words ("ChatGPT" -> "Chat" + "G" + "PT").

Byte-Pair Encoding (BPE)

strawberry -> [“str”, “aw”, “berry”]

Byte-Pair Encoding (BPE)

“London”

“Londres”

“Londra”

Byte-Pair Encoding (BPE)

“London”	(L, o) 3
“Londres”	(o, n) 4
“Londra”	(n, d) 3
{L, o, n, d, r, e, s, a}	(d, o) 1
	(d, r) 2
	(r, e) 1
	(e, s) 1
	(r, a) 1

Byte-Pair Encoding (BPE)

“London”	(L, o) 3
“Londres”	(o, n) 4
“Londra”	(n, d) 3
{L, o, n, d, r, e, s, a, on}	(d, o) 1
rule #1: (o, n) -> on	(d, r) 2
	(r, e) 1
	(e, s) 1
	(r, a) 1

Byte-Pair Encoding (BPE)

“L <u>on</u> d <u>on</u> ”	(L, on) 3
“L <u>on</u> d r e s”	(on, d) 3
“L <u>on</u> d r a”	(d, on) 1
	(d, r) 2
{L, o, n, d, r, e, s, a, on}	(r, e) 1
rule #1: (o, n) -> on	(e, s) 1
	(r, a) 1

Byte-Pair Encoding (BPE)

“L on d on”

(L, on) 3

“L on d r e s”

(on, d) 3

“L on d r a”

(d, on) 1

(d, r) 2

{L, o, n, d, r, e, s, a, on, **Lon**}

(r, e) 1

(e, s) 1

rule #1: (o, n) -> on

rule #2: (L, on) -> Lon

(r, a) 1

Byte-Pair Encoding (BPE)

“Lon d on” (Lon, d) 3

“Lon d r e s ” (d, on) 1

“Lon d r a” (d, r) 2

{L, o, n, d, r, e, s, a, on, Lon} (r, e) 1

(e, s) 1

rule #1: (o, n) -> on

rule #2: (L, on) -> Lon (r, a) 1

Byte-Pair Encoding (BPE)

“Lon d on”

(Lon, d) 3

“Lon d r e s”

(d, on) 1

“Lon d r a”

(d, r) 2

{L, o, n, d, r, e, s, a, on, Lon, **Lond**}

(r, e) 1

rule #1: (o, n) -> on

(e, s) 1

rule #2: (L, on) -> Lon

(r, a) 1

rule #3: (Lon, d) -> Lond

Byte-Pair Encoding (BPE)

“Lond on”

(Lond, on) 1

“Lond r e s ”

(Lond, r) 2

“Lond r a”

(r, e) 1

{L, o, n, d, r, e, s, a, on, Lon, Lond}

(e, s) 1

rule #1: (o, n) -> on

(r, a) 1

rule #2: (L, on) -> Lon

rule #3: (Lon, d) -> Lond

Byte-Pair Encoding (BPE)

“Lond on”

(Lond, on) 1

“Lond r e s ”

(Lond, r) 2

“Lond r a”

(r, e) 1

{L, o, n, d, r, e, s, a, on, Lon, Lond,
Londr}

(e, s) 1

rule #1: (o, n) -> on

rule #2: (L, on) -> Lon

(r, a) 1

rule #3: (Lon, d) -> Lond

rule #4: (Lond, r) -> Londr

Byte-Pair Encoding (BPE)

{L, o, n, d, r, e, s, a, on, Lon,
Lond, Londr}

Vocabulary

rule #1: (o, n) -> on

rule #2: (L, on) -> Lon

rule #3: (Lon, d) -> Lond

rule #4: (Lond, r) -> Londr

Merge Rules

Byte-Pair Encoding (BPE)

{L, o, n, d, r, e, s, a, on, Lon,
Lond, Londr}

rule #1: (o, n) -> on

rondLon

rule #2: (L, on) -> Lon

rule #3: (Lon, d) -> Lond

rule #4: (Lond, r) -> Londr

Byte-Pair Encoding (BPE)

{L, o, n, d, r, e, s, a, on, Lon,
Lond, Londr}

rule #1: (o, n) -> on

r on d L on

rule #2: (L, on) -> Lon

rule #3: (Lon, d) -> Lond

rule #4: (Lond, r) -> Londr

Byte-Pair Encoding (BPE)

{L, o, n, d, r, e, s, a, on, Lon,
Lond, Londr}

rule #1: (o, n) -> on

r on d L on

rule #2: (L, on) -> Lon

rule #3: (Lon, d) -> Lond

rule #4: (Lond, r) -> Londr

Byte-Pair Encoding (BPE)

{L, o, n, d, r, e, s, a, on, Lon,
Lond, Londr}

rule #1: (o, n) -> on

r on d Lon

rule #2: (L, on) -> Lon

rule #3: (Lon, d) -> Lond

rule #4: (Lond, r) -> Londr

Byte-Pair Encoding (BPE)

{L, o, n, d, r, e, s, a, on, Lon,
Lond, Londr}

rule #1: (o, n) -> on

r on d Lon

rule #2: (L, on) -> Lon

rule #3: (Lon, d) -> Lond

No merges for #3 and #4

rule #4: (Lond, r) -> Londr

Byte-Pair Encoding (BPE)

{L, o, n, d, r, e, s, a, on, Lon,
Lond, Londr}

rule #1: (o, n) -> on

rondLon

rule #2: (L, on) -> Lon

r on d Lon

rule #3: (Lon, d) -> Lond

rule #4: (Lond, r) -> Londr

Byte-Pair Encoding (BPE)

{0: L, 1: o, 2:n, 3: d,
4: r, 5: e, 6: s,
7: a, 8: on, 9: Lon,
10: Lond, 11: Londr}

rondLon

r on d Lon

[4, 8, 3, 9]

Why is BPE important?

“London”	[23421]	[4222]
“Londres”	[43, 623, 411]	[26432]
“Londra”	[43, 623, 430]	[17681, 520]

GPT-2

Mistral

BPE is data sensitive

90% general web

5% programming

5% math

st raw berry

80% general web

10% programming

10% math

str aw berry

BPE is data sensitive

90% general web

5% programming

5% math

+ emoji

build tokenizer

general web

programming

math

(remove all emojis)

pretraining data

BPE is data sensitive

Fishing for Magikarp: Automatically Detecting Under-trained Tokens in Large Language Models

Thank you!